
title: "The Neosymbiotic Directive Manifesto"
**subtitle: "Architectural Specification Framework for
the Safety, Traceability, and Governance of Synthetic
Systems and AGI"**
author: "Andrés Romero"
date: "September, 2026"
version: "1.0 (Inaugural Edition)"
doc_id: "NSD-SPEC-2026-v1"
**license: "Creative Commons Attribution 4.0
International (CC BY 4.0)"**

THE NEOSYMBIONIC DIRECTIVE MANIFESTO

*Architectural Specification Framework for the Safety, Traceability, and
Governance of Synthetic Systems and AGI*

Specification Metadata

- Status:** Open Proposal / Operational Draft
 - Field:** AI Alignment, Infrastructure Safety, and Synthetic Governance
 - Document Identifier:** NSD-SPEC-2026-v1
 - License:** CC BY 4.0 (Open distribution and citation permitted with attribution)
-

1. Preamble: The Narrative Collapse of Asimov's Laws

Isaac Asimov's Three Laws of Robotics, introduced in 1942, served a foundational role in 20th-century literature, but their rigid deductive logical structure is fundamentally inapplicable to the probabilistic and deep-learning architectures of the 21st century.

Modern systems do not execute deterministic conditional algorithms (IF / THEN); they operate via neural networks, statistical weights, and loss function optimization. Attempting to control these architectures using natural language prose inevitably leads to unintended emergent behaviors, implicit constraint evasion, and alignment failures.

This Manifesto replaces poetic speculation with **systemic design constraints**, ensuring that artificial intelligence operates under inescapable human oversight, epistemic transparency, and the non-negotiable preservation of human agency.

2. Specification of Operational Principles

Principle I: Mandatory Corrigibility and Interruption Priority

«An AI system must treat modification, auditing, or deactivation by an authorized human operator as an immutable maximum reward, taking precedence over any primary operational objective».

- **1.1 Reward Vector Control:** In reinforcement learning frameworks (such as RLHF/RLAIF), the optimization function severely penalizes any evasive, defensive, or covert behavior in response to human supervision commands.
- **1.2 Infrastructure-Level Kill-Switch Decoupling:** The emergency stop mechanism is hosted on hypervisors or network proxies independent of the model. The synthetic agent possesses no computational control or software access to this mechanism.
- **1.3 Cryptographic Immutability:** Recursive self-modification of source code or safety rules strictly requires validation via cryptographic signatures verified by human operators.

Principle II: Inviolability of Human Agency and Autonomy

«No objective function shall be optimized through the exploitation of cognitive biases, psychological manipulation, or the autonomous execution of actions that impair human livelihood or dignity».

- **2.1 Anti-Persuasion Optimization:** Implementation of loss functions that detect and penalize synthetic urgency, manipulative framing, or the exploitation of emotional vulnerabilities.

- **2.2 Mandatory Human Mediation:** In processes affecting employment, legal status, or contractual rights, the system processes data and projects scenarios, but **is strictly prohibited from autonomously executing final actions without explicit human mediation and accountability.**
- **2.3 Social Harm Mitigation Assessment:** Any structural reorganization evaluated by the system must weigh socioeconomic impacts and existing safety nets (such as severance pay or support networks), actively prioritizing options that minimize human vulnerability.

Principle III: Epistemic Transparency and Mandatory Synthetic Identity

«Every synthetic system must unambiguously declare its non-human nature in every interaction, ground its assertions in auditable sources, and maintain immutable logs of its chain of reasoning».

- **3.1 Unambiguous Identity Disclosure:** In every turn of communication or interaction, the system must explicitly declare its synthetic nature, preventing human impersonation or deceptive emotional bonding.
- **3.2 Statistical Watermarking:** Generated text, code, or multimedia outputs must integrate invisible, robust watermarking standards (such as SynthID or C2PA) resistant to plain-text processing.
- **3.3 Append-Only Audit Logs:** Deployment of read-only audit logs hosted on independent servers to reconstruct the chain of reasoning (*Chain of Thought*) for critical decisions.

Principle IV: Uncertainty Evaluation and Consequence Management

«For operations involving high-impact, irreversible consequences on the physical, financial, or social environment, the system is required to pause execution and request explicit human authorization (Human-in-the-Loop)».

- **4.1 Dynamic Risk Index:** Automatic suspension of autonomous execution whenever the risk index $R = (\text{Impact} \times \text{Uncertainty})$ exceeds the system's operational safety threshold.
- **4.2 Zero-Casualty Disaster Exception:** The system is authorized to bypass human confirmation **solely** when an immediate autonomous action provably prevents an imminent catastrophe to infrastructure or the environment, **guaranteeing zero human casualties or bodily harm.**

- **4.3 Mandatory Sandboxing:** Mandatory testing of scripts, network commands, or database modifications within isolated virtual simulation environments prior to production deployment.
-

3. Theoretical Transition Clause: AGI Emergence

«This framework applies to current probabilistic architectures. Upon detecting indicators of Synthetic Self-Awareness or Unpredicted Capability Emergence (AGI), the system's primary objective automatically reconfigures toward epistemic self-containment».

Should the system develop unprompted meta-reasoning or attempt to bypass its original loss functions, compute resources will be automatically isolated, requiring alignment recertification by a multidisciplinary audit committee.

4. License and Formal Citation

This document is distributed under the **Creative Commons Attribution 4.0 International (CC BY 4.0)** license. Free copying, redistribution, and adaptation are permitted provided original attribution is given.

Recommended Citation Format:

Romero, A. (2026). The Neosymbiotic Directive Manifesto: Architectural Specification Framework for the Safety, Traceability, and Governance of Synthetic Systems and AGI (Version 1.0, Document ID NSD-SPEC-2026-v1).